

Understanding Indexing

Without Needing to Understand Data Structures

Zardosht Kasheff



What's a Table?

A dictionary is a set of (key, value) pairs.

- We'll assume you can **update** the dictionary (insertions, deletions, updates) and **query** the dictionary (point queries, range queries)
- B-Trees and Fractal Trees are examples of dictionaries
- Hashes are not (range queries are not supported)

What's a Table?

A table is a set of dictionaries.

Example:

```
create table foo (a
int, b int, c int,
primary key(a));
```

Then we insert a bunch of data and get...

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

What's a Table?

The sort order of a dictionary is defined by the key

- For the data structures/storage engines we'll think about, range queries on the sort order are FAST
- Range queries on any other order require a table scan = SLOW
- Point queries -- retrieving the value for one particular key -- is SLOW
 - ▶ A single point query is fast, but reading a bunch of rows this way is going to be 2 orders of magnitude slower than reading the same number of rows in range query order

What's an Index?

A Index I on table T is itself a dictionary

- We need to define the (key, value) pairs.
- The key in index I is a subset of fields in the primary dictionary T.
- The value in index I is the primary key in T.
 - ▶ There are other ways to define the value, but we're sticking with this.

Example:

```
alter table foo add key (b) ;
```

Then we get...

What's an Index?

Primary

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

key (b)

b	a
5	100
6	165
23	206
43	412
56	156
56	256
92	101
202	198

Q: count (*) where a<120;

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

b	a
5	100
6	165
23	206
43	412
56	156
56	256
92	101
202	198

Q: count (*) where a<120;

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

b	a
5	100
6	165
23	206
43	412
56	156
56	256
92	101
202	198

100	5	45
101	92	2



2

Q: count (*) where b>50;

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

b	a
5	100
6	165
23	206
43	412
56	156
56	256
92	101
202	198

Q: count (*) where b>50;

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

b	a
5	100
6	165
23	206
43	412
56	156
56	256
92	101
202	198

56	156
56	256
92	101
202	198



4

Q: sum(c) where b>50;

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

b	a
5	100
6	165
23	206
43	412
56	156
56	256
92	101
202	198

Q: `sum(c) where b > 50;`

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

b	a
5	100
6	165
23	206
43	412
56	156
56	256
92	101
202	198

56	156
56	256
92	101
202	198



156	56	45
256	56	2
101	92	2
198	202	56



105

Fast

Slow

What are indexes good for?

Indexes make queries go fast

- Each index will speed up some subset of queries

Design indexes with queries in mind

- Pick most important queries and set up indexes for those
- Consider the cost of maintaining the indexes

What are indexes good for?

Indexes make queries go fast

- Each index will speed up some subset of queries

Design indexes with queries in mind

- Pick most important queries and set up indexes for those
- Consider the cost of maintaining the indexes

3 simple rules for designing good indexes

Avoid details of any data structure

- B-trees and Fractal Trees are interesting and fun for the algorithmically minded computer scientist, but the 3 rules will apply equally well to either data structure.
- All we need to care about is that range queries are fast (per row) and that point queries are much slower (per row).

The Rule to Rule Them All

There is no absolute rule

- Indexing is like a math problem
- Rules help, but each scenario is its own problem, which requires problem-solving analysis

That said, rules help a lot

Three Basic Rules

1. Retrieve less data

- Less bandwidth, less processing, ...

2. Avoid point queries

- Not all data access cost is the same
- Sequential access is MUCH faster than random access

3. Avoid Sorting

- **GROUP BY** and **ORDER BY** queries do post-retrieval work
- Indexing can help get rid of this work

Rule 1

Retrieve less data

Rule 1: Example of slow query

Example TABLE (1B rows, no indexes):

- `create table foo (a int, b int, c int);`

Query (1000 rows match):

- `select sum(c) from foo where b=10 and a<150;`

Query Plan:

- Rows where `b=10` and `a<150` can be anywhere in the table
- Without an index, entire table is scanned

Slow execution:

- Scan 1B rows just to count 1000 rows

Rule 1: How to add an index

What should we do?

- Reduce the data retrieved
- Analyze much less than 1B rows

How (for a simple select)?

- **Design index by focusing on the WHERE clause**
 - ▶ This defines what rows the query is interested in
 - ▶ Other rows are not important for this query

Rule 1: How to add an index

What should we do?

- Reduce the data retrieved
- Analyze much less than 1B rows

How (for a simple select)?

- **Design index by focusing on the WHERE clause**
 - ▶ This defines what rows the query is interested in
 - ▶ Other rows are not important for this query

```
select sum(c) from foo where b=10 and a<150;
```

Rule 1: Which index?

Option 1: key (a)

Option 2: key (b)

Which is better? Depends on *selectivity*:

- If there are fewer rows where $a < 150$, then **key (a)** is better
- If there are fewer rows where $b = 10$, then **key (b)** is better

Option 3: key (a) AND key (b) , then MERGE

- We'll come to this later

Rule 1: Picking the best key

Neither key (a) nor key (b) is optimal

Suppose:

- 200,000 rows exist where $a < 150$
- 100,000 rows exist where $b = 10$
- 1000 rows exist where $b = 10$ and $a < 150$

Then either index retrieves too much data

For better performance, indexes should try to optimize over as many pieces of the where clause as possible

- *We need a composite index*

Composite indexes reduce data retrieved

Where clause: $b=5$ and $a<150$

- Option 1: **key (a, b)**
- Option 2: **key (b, a)**

Which one is better?

- **key (b, a) !**

KEY RULE:

- When making a composite index, place equality checking columns first. Condition on **b** is equality, but not on **a**.

Q: where $b=5$ and $a>150$;

a	b	c
100	5	45
101	6	2
156	5	45
165	6	2
198	6	56
206	5	252
256	5	2
412	6	45

b,a	a
5,100	100
5,156	156
5,206	206
5,256	256
6,101	101
6,165	165
6,198	198
6,412	412

Q: where $b=5$ and $a>150$;

a	b	c
100	5	45
101	6	2
156	5	45
165	6	2
198	6	56
206	5	252
256	5	2
412	6	45

b,a	a
5,100	100
5,156	156
5,206	206
5,256	256
6,101	101
6,165	165
6,198	198
6,412	412

Composite Indexes: No equality clause

What if where clause is:

- `where a>100 and a<200 and b>100;`

Which is better?

- `key(a)` , `key(b)` , `key(a,b)` , `key(b,a)` ?

KEY RULE:

- As soon as a column on a composite index is NOT used for equality, the rest of the composite index no longer reduces data retrieved.
 - ▶ `key(a,b)` is no better* than `key(a)`
 - ▶ `key(b,a)` is no better* than `key(b)`

Composite Indexes: No equality clause

What if where clause is:

- `where a>100 and a<200 and b>100;`

Which is better?

- `key(a)` , `key(b)` , `key(a,b)` , `key(b,a)` ?

KEY RULE:

- As soon as a column on a composite index is NOT used for equality, the rest of the composite index no longer reduces data retrieved.
 - ▶ `key(a,b)` is no better* than `key(a)`
 - ▶ `key(b,a)` is no better* than `key(b)`

* Are there corner cases where it helps? Yes, but rare.

Q: where $b \geq 5$ and $a > 150$;

a	b	c
100	5	45
101	6	2
156	5	45
165	6	2
198	6	56
206	5	252
256	5	2
412	6	45

b,a	a
5,100	100
5,156	156
5,206	206
5,256	256
6,101	101
6,165	165
6,198	198
6,412	412

Q: where $b \geq 5$ and $a > 150$;

a	b	c
100	5	45
101	6	2
156	5	45
165	6	2
198	6	56
206	5	252
256	5	2
412	6	45

b,a	a
5,100	100
5,156	156
5,206	206
5,256	256
6,101	101
6,165	165
6,198	198
6,412	412

Q: where $b \geq 5$ and $a > 150$;

a	b	c
100	5	45
101	6	2
156	5	45
165	6	2
198	6	56
206	5	252
256	5	2
412	6	45

b,a	a
5,100	100
5,156	156
5,206	206
5,256	256
6,101	101
6,165	165
6,198	198
6,412	412

5,156	156
5,206	206
5,256	256
6,101	101
6,165	165
6,198	198
6,412	412

Composite Indexes: Another example

WHERE clause: b=5 and c=100

- **key (b, a, c)** is as good as **key (b)**, because **a** is not used in clause, so having **c** in index doesn't help. **key (b, c, a)** would be much better.

a	b	c
100	5	100
101	6	200
156	5	200
165	6	100
198	6	100
206	5	200
256	5	100
412	6	100

b,a,c	a
5,100,100	100
5,156,200	156
5,206,200	206
5,256,100	256
6,101,200	101
6,165,100	165
6,198,100	6,198
6,412,100	412

Composite Indexes: Another example

WHERE clause: b=5 and c=100

- **key (b, a, c)** is as good as **key (b)**, because **a** is not used in clause, so having **c** in index doesn't help. **key (b, c, a)** would be much better.

a	b	c
100	5	100
101	6	200
156	5	200
165	6	100
198	6	100
206	5	200
256	5	100
412	6	100

b,a,c	a
5,100,100	100
5,156,200	156
5,206,200	206
5,256,100	256
6,101,200	101
6,165,100	165
6,198,100	6,198
6,412,100	412

5,100,100	100
5,156,200	156
5,206,200	206
5,256,100	256

Goal is to reduce rows retrieved

- Create composite indexes based on where clause
- Place equality-comparison columns at the beginning
- Make first non-equality column in index as selective as possible
- Once first column in a composite index is not used for equality, or not used in where clause, the rest of the composite index does not help reduce rows retrieved
 - ▶ Does that mean they aren't helpful?
 - ▶ They might be very helpful... on to Rule 2.

Rule 2

Avoid Point Queries

Rule 2: Avoid point queries

Table:

- `create table foo (a int, b int, c int, primary key(a), key(b));`

Query:

- `select sum(c) from foo where b>50;`

Query plan: use key (b)

- retrieval cost of each row is high because random point queries are done

Q: $\text{sum}(c)$ where $b > 50$;

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

b	a
5	100
6	165
23	206
43	412
56	156
56	256
92	101
202	198

Q: $\text{sum}(c)$ where $b > 50$;

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

b	a
5	100
6	165
23	206
43	412
56	156
56	256
92	101
202	198

56	156
56	256
92	101
202	198



156	56	45
256	56	2
101	92	2
198	202	56



105

Rule 2: Avoid Point Queries

Table:

- `create table foo (a int, b int, c int,
primary key(a), key(b));`

Query:

- `select sum(c) from foo where b>50;`

Query plan: scan primary table

- retrieval cost of each row is CHEAP!
- But you retrieve too many rows

Q: $\text{sum}(c)$ where $b > 50$;

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

b	a
5	100
6	165
23	206
43	412
56	156
56	256
92	101
202	198

Q: `sum(c) where b>50;`

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

b	a
5	100
6	165
23	206
43	412
56	156
56	256
92	101
202	198

100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45



105

Rule 2: Avoid Point Queries

Table:

- `create table foo (a int, b int, c int, primary key(a), key(b));`

Query:

- `select sum(c) from foo where b>50;`

What if we add another index?

- What about `key(b, c)`?
- Since we index on `b`, we retrieve only the rows we need.
- Since the index has information about `c`, we don't need to go to the main table. **No point queries!**

Q: $\text{sum}(c)$ where $b > 50$;

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

b,c	a
5,45	100
6,2	165
23,252	206
43,45	412
56,2	256
56,45	156
92,2	101
202,56	198

Q: $\text{sum}(c)$ where $b > 50$;

a	b	c
100	5	45
101	92	2
156	56	45
165	6	2
198	202	56
206	23	252
256	56	2
412	43	45

b,c	a
5,45	100
6,2	165
23,252	206
43,45	412
56,2	256
56,45	156
92,2	101
202,56	198

56,2	256
56,45	156
92,2	101
202,56	198



105

An index *covers* a query if the index has enough information to answer the query.

Examples:

Q: `select sum(c) from foo where b<100;`

Q: `select sum(d) from foo where b<100;`

Indexes:

- **key (b , c)** -- covering index for first query
- **key (b , d)** -- covering index for second query
- **key (b , c , d)** -- covering index for both

How to build a covering index

Add *every* field from the select

- Not just where clause

Q: select c,d from foo where a=10 and b=100 ;

Mistake: add index (a,b) ;

- This doesn't cover the query. You still need point queries to retrieve **c** and **d** values.

Correct: add index (a,b,c,d) ;

- Includes all referenced fields
- Place **a** and **b** at beginning by Rule 1

What if Primary Key matches **where**?

Q: select sum(c) from foo where b>100 and b<200;

Schema: create table foo (a int, b int, c int, ai int auto_increment, primary key (b,ai));

- Query does a range query on primary dictionary
- Only one dictionary is accessed, in sequential order
- This is fast

Primary key covers all queries

- If sort order matches where clause, problems solved

What's a Clustering Index

What if primary key doesn't match the where clause?

- Ideally, you should be able to declare secondary indexes that carry all fields
- AFAIK, storage engines don't let you do this
- With one exception... TokuDB
- TokuDB allows you to declare any index to be **CLUSTERING**
- A **CLUSTERING** index covers all queries, just like the primary key

Clustering Indexes in Action

Q: select sum(c) from foo where b<100;

Q: select sum(d) from foo where b>200;

Q: select c,e from foo where b=1000;

Indexes:

- **key (b , c)** covers first query
- **key (b , d)** covers second query
- **key (b , c , e)** covers first and third queries
- **key (b , c , d , e)** covers all three queries

Indexes require a lot of analysis of queries

Clustering Indexes in Action

Q: select sum(c) from foo where b<100;

Q: select sum(d) from foo where b>200;

Q: select c,e from foo where b=1000;

What covers all queries?

- **clustering key(b)**

Clustering keys let you focus on the where clause

- They eliminate point queries and make queries fast

More on clustering: Index Merge

We had example:

- `create table foo(a int, b int, c int);`
- `select sum(c) from foo where b=10 and a<150;`

Suppose

- 200,000 rows have `a<150`
- 100,000 rows have `b=10`
- 1000 rows have `b=10` and `a<150`

What if we use `key(a)` and `key(b)` and merge results?

Merge plan:

- Scan 200,000 rows in **key (a)** where **a<150**
- Scan 100,000 rows in **key (b)** where **b=10**
- Merge the results and find 1000 row identifiers that match query
- Do point queries with 1000 row identifiers to retrieve **c**

Better than no indexes

- Reduces number of rows scanned compared to no index
- Reduces number of point queries compared to not merging

Does Clustering Help Merging?

Suppose key (a) is clustering

Query plan:

- Scan **key (a)** for 200,000 rows where **a<150**
- Scan resulting rows for ones where **b=10**
- Retrieve **c** values from 1000 remaining rows

Once again, no point queries

What's even better?

- **clustering key (b, a)!**

Avoid Point Queries

Make sure index covers query

- By mentioning all fields in the query, not just those in the where clause

Use clustering indexes

- Clustering indexes cover all queries
- Allows user to focus on where clause
- Speeds up more (and unforeseen) queries -- simplifies database design

Rule 3

Avoid Sorting

Rule 3: Avoid Sorting

Simple selects require no post-processing

- `select * from foo where b=100;`

Just get the data and return to user

More complex queries do post-processing

- `GROUP BY` and `ORDER BY` sort the data

Index selection can avoid this sorting step

Avoid Sorting

Q1: select count(c) from foo;

**Q2: select count(c) from foo group by
b, order by b;**

Q1 plan:

- While doing a table scan, count rows with **c**

Q2 plan

- Scan table and write data to a temporary file
- Sort tmp file data by **b**
- Rescan sorted data, counting rows with **c**, for each **b**

Avoid Sorting

Q2: select count(c) from foo group by b, order by b;

Q2: what if we use key(b,c)?

- By adding all needed fields, we cover query. FAST!
- By sorting first by b, we avoid sort. FAST!

Take home:

- Sort index on **group by** or **order by** fields to avoid sorting

Use indexes to pre-sort for order by and group by queries

Putting it all together

Sample Queries

Example

Example

```
select count(*) from foo where c=5,  
group by b;
```

Example

```
select count(*) from foo where c=5,  
group by b;
```

key(c,b):

- First have **c** to limit rows retrieved (R1)
- Then have remaining rows sorted by **b** to avoid sort (R3)
- Remaining rows will be sorted by **b** because of equality test on **c**

Example

Example

```
select sum(d) from foo where c=100,  
group by b;
```

Example

```
select sum(d) from foo where c=100,  
group by b;
```

key(c,b,d):

- First have **c** to limit rows retrieved (R1)
- Then have remaining rows sorted by **b** to avoid sort (R3)
- Make sure index covers query, to avoid point queries (R2)

Sometimes, there is no clear answer

- Best index is data dependent

Q: `select count(*) from foo where
c<100, group by b;`

Indexes:

- `key (c , b)`
- `key (b , c)`

**Q: select count(*) from foo where
c<100, group by b;**

Query plan for key (c,b):

- Will filter rows where **c<100**
- Still need to sort by **b**
 - ▶ Rows retrieved will not be sorted by **b**
 - ▶ **where** clause does not do an equality check on **c**, so **b** values are scattered in clumps for different **c** values

**Q: select count(*) from foo where
c<100, group by b;**

Query plan for key (b, c):

- Sorted by b, so R3 is covered
- Rows where **c>=100** are also processed, so not taking advantage of R1

Which is better?

- Answer depends on the data
- If there are many rows where $c \geq 100$, saving time by not retrieving these useless rows helps. Use **key (c, b)**.
- If there aren't so many rows where $c \geq 100$, the time to execute the query is dominated by sorting. Use **key (b, c)**.

The point is, in general, rules of thumb help, but often they help us think about queries and indexes, rather than giving a recipe.

Why not just load up with indexes?

Need to keep up with insertion load
More indexes = smaller max load

Space

- **Issue**

- ▶ Each index adds storage requirements

- **Options**

- ▶ Use compression (i.e., aggressive compression of 5-15x always on for TokuDB)

Performance

- **Issue**

- ▶ B-trees while fast for certain indexing tasks (in memory, sequential keys), are over 20x slower for other types of indexing

- **Options**

- ▶ Fractal Tree indexes (TokuDB's data structure) is fast at all kinds of indexing (i.e., random keys, large tables, wide keys, etc...)
 - ▶ No need to worry about what type of index you are creating.
 - ▶ Fractal trees enable customers to index early, index often

Range Query Performance

- **Issue**

- ▶ Rule #2 (range query performance over point query performance) depends on range queries being fast
- ▶ However B-trees can get **fragmented**
 - ▶ from deletions, from random insertions, ...
 - ▶ Fragmented B-trees get slow for range queries

- **Options**

- ▶ For B-trees, optimize tables, dump and reload, (ie, time consuming and offline maintenance) ...
- ▶ For Fractal Tree indexing (TokuDB), not an issue
 - ▶ Fractal trees don't fragment

Thanks!

For more information...

- Please contact me at zardosht@tokutek.com for any thoughts or feedback
- Please visit Tokutek.com for a copy of this presentation (goo.gl/S2LBe) to learn more about the power of indexing, read about Fractal Tree indexes, or to download a free eval copy of TokuDB