




Java. Cloud. Leadership.

Transactions: Over used or just misunderstood?

Dr Mark Little
Red Hat, Inc.
25/9/2012

ACID Properties

- Atomicity
- *Consistency*
- Isolation
- Durability




Atomicity: two-phase commit

- Required when there is more than one resource managers (RM) in a transaction
- Managed by the transaction manager (TM)
- Uses a familiar two-phase technique:



Consensus is not enough

- Coordinator must:
 - maintain a transaction log
 - be able to communicate the resolution strategy to each of the participants
- During recovery, coordinator may need to contact resource managers directly
 - Some variations of protocol require resource managers to contact coordinator




4



Java. Cloud. Leadership.

Isolation

- Programs can be executed concurrently
 - BUT they appear to execute serially
- P1 : $z := 10; x := x+1; y := y+1$
- P2 : $w := 7; x := x * 2; y := y * 2$
- $(z := 10 \parallel w := 7); x := x+1; y := y+1; x := x * 2; y := y * 2$
- $(z := 10 \parallel w := 7); x := x+1; x := x * 2; y := y * 2; y := y + 1$



5



Java. Cloud. Leadership.

Durability

- When a transaction commits, its results must survive failures
 - Must be durably recorded prior to commit
 - System waits for disk ack before acking user
- If a transaction rolls back, changes are undone
- Durability is probabilistic
 - It does not have to be disk based!



6



Java. Cloud. Leadership.

Heuristics

- Two-phase commit protocol is blocking
- To break the blocking nature, prepared participants may make autonomous decisions to commit or rollback
 - Participant must durably record this decision
 - If the decision differs then possibly non-atomic outcome
- Heuristics cannot be resolved automatically




7



Java. Cloud. Leadership.

Optimisations


- one phase commit
 - no voting if transaction tree is single branch



One Phase Commit

- “read-only”
 - ✓ resource doesn't change any data
 - ✓ can be ignored in second phase of commit

Interposition



“They add overhead for no benefit”

- Many customers regularly run with multiple one-phase only participants in the same transaction
 - And appear to get atomicity and failure recovery
 - Failures don't happen that often
- Lack of education on the problem
 - Often surprised when risks are explained
 - But still want atomicity and failure recovery with no impact on performance



10



Java. Cloud. Leadership.

“I don't always need atomicity”

- Forward compensation transactions may offer a medium-term solution
 - Do-undo
- The basis for all of the Web Services transactions specifications and standards
- May provide a viable upgrade path for multiple 1PC participants too
 - Not quite ACID, but better than manual/ad hoc

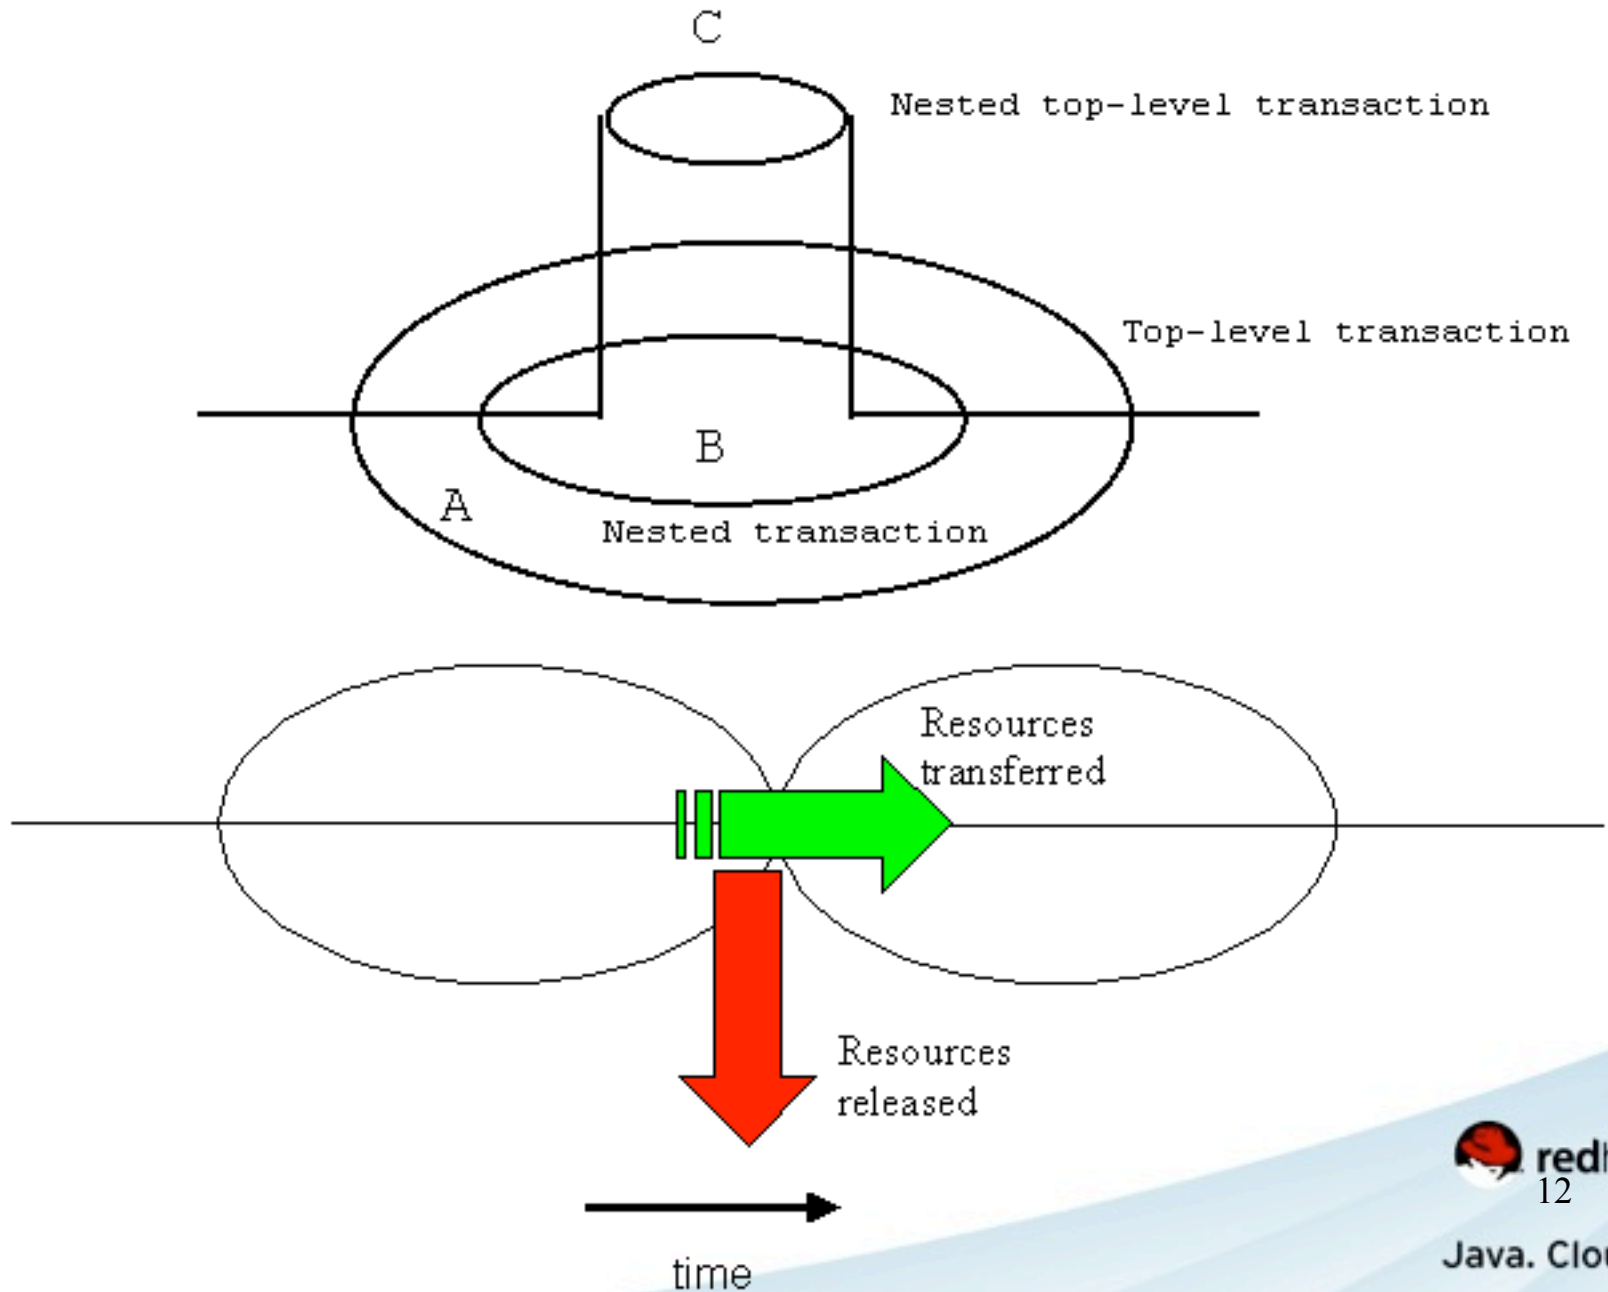


11



Java. Cloud. Leadership.

Relaxing ACID



When to use transactions

- When you need ACID semantics!
- Or ...
 - When you have a need to guarantee consensus in the presence of failures
 - When you need isolation and consistency across failures
- Relaxing ACID semantics is possible
- Recoverable transactions may be sufficient

When not to use transactions

- When all you want is consensus
- When you will only ever have a single resource
 - Though 1PC optimisations can make overhead negligible



14



Java. Cloud. Leadership.

When not to cut corners

- If you want ACID semantics then use an implementation that provides them all
- If you want to relax the semantics then use an implementation that allows that to happen
- Don't use an “ACID” transaction system that doesn't provide for ACID
- Don't expect ACID guarantees when your application doesn't uphold its end of the contract



15



Java. Cloud. Leadership.

Conclusions

- Atomicity requires durability if failures are to be survived
- Failures require recovery
- Two-phase commit not just for distributed cases
- Extended transactions useful for new data strategies
- Many cases of 1 resource becoming 2 or more
 - Transactions take care of that transition opaquely